

# Decoupling Workflow Services

14 March 2024 14:44

1. Tight Coupling vs. Loose Coupling:
  - Tight coupling between components in an application can lead to single points of failure and scalability issues.
  - Loose coupling involves decoupling components, allowing them to operate independently and communicate through managed services like load balancers or messaging queues.
  - Loose coupling ensures resilience and scalability by removing direct dependencies between components.
2. Services for Decoupling:
  - Simple Queue Service (SQS): A fully-managed, highly-available messaging tool for decoupling applications and managing message queues.
  - Simple Notification Service (SNS): Used for pushing out notifications to endpoints, such as applications, text messages, or emails.
  - API Gateway: Provides a scalable and secure front door to applications, controlling access to resources and enabling communication with backend services.
3. Exam Tips:
  - Always prioritize loose coupling over tight coupling in architectural designs.
  - Avoid direct EC2-to-EC2 calls and instead use managed services like load balancers or messaging queues.
  - Understand the use cases and benefits of SQS, SNS, and API Gateway for decoupling applications.
4. Additional Tips:
  - Remember that there's no one-size-fits-all solution for decoupling; choose the appropriate service based on the specific requirements of your application.
  - Stay calm and think strategically, even in challenging situations, like a bull charging at you (metaphorically speaking).

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

5. Poll-Based Messaging:
  - Poll-based messaging involves a producer (sender) placing messages in a queue, and consumers (receivers) periodically polling the queue for messages.
  - SQS acts as a messaging queue, enabling asynchronous processing of workloads by decoupling message producers and consumers.
6. SQS Settings:
  - Delivery Delay: Delays the delivery of messages to consumers for a specified period, allowing time for additional processing or validation.
  - Message Size: Messages can be up to 256 kilobytes in size, regardless of format.
  - Encryption: Messages can be encrypted in transit and at rest using AWS Key Management Service (KMS) for enhanced security.
  - Message Retention: Messages are retained in the queue for a specified period (default is 4 days, maximum is 14 days).
  - Polling: SQS supports short polling (default) and long polling. Long polling reduces the number of API calls and CPU usage.
  - Queue Depth: CloudWatch alarms can be triggered based on the depth of the queue, enabling auto-scaling based on workload.
7. Visibility Timeout:
  - The visibility timeout is the duration during which a message is invisible to other consumers after being received by a consumer.
  - If the message is not deleted or processed within the visibility timeout period, it becomes visible again in the queue for other consumers.

## 8. Exam Tips:

- Understand all SQS settings and their implications, as they are commonly tested in the exam.
- Remember the limitations of SQS, such as message retention and maximum message size.
- Know the difference between short polling and long polling, and when to use each method.
- Be prepared to troubleshoot SQS issues, such as lost messages or unexpected message visibility.
- Stay attentive to details, as questions may involve scenarios where misconfigured settings lead to unexpected behavior.

## 9. Additional Tips:

- Remember to spread food evenly on a plate when microwaving to ensure even heating and avoid cold spots.

Overall, mastering SQS and its settings is crucial for designing scalable and reliable distributed systems on AWS. Make sure to review these concepts thoroughly, and stay mindful of their practical applications in real-world scenarios.

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

## 1. What are Dead-Letter Queues?:

- Dead-letter queues are target queues for messages that cannot be processed successfully from the primary message queue.
- They serve as a failover mechanism to store messages for inspection and potential replay later.

## 2. Features of Dead-Letter Queues:

- Dead-letter queues work with both SQS and Amazon SNS.
- They are useful for debugging applications and messaging systems, allowing isolation of unconsumed messages for troubleshooting.
- Messages can be redriven from the dead-letter queue back to the original source queue.

## 3. Configuration and Benefits:

- Dead-letter queues are regular SQS queues that are configured to receive failed messages.
- Benefits include configuring alarms based on message availability counts, identifying specific log files for investigation, analyzing message contents for errors, and troubleshooting consumer permissions.

## 4. Visualizing Dead-Letter Queue Workflow:

- A visual representation demonstrates how dead-letter queues work within a messaging-based application system, including the transfer of failed messages to the dead-letter queue for further analysis and potential replay.

## 5. Hands-On Demonstration:

- The lesson includes a hands-on demonstration in the AWS Sandbox environment, setting up SQS queues, configuring dead-letter queues, and testing message processing with Lambda functions.

## 6. Exam Tips:

- Important exam tips include monitoring queue depth with CloudWatch alarms, understanding that dead-letter queues are regular SQS queues, ensuring compatibility with FIFO queues if required, and utilizing dead-letter queues with Amazon SNS topics.

Overall, understanding dead-letter queues is essential for effectively managing message processing and troubleshooting in distributed systems on AWS. Make sure to review these concepts thoroughly, as they are likely to be tested in AWS certification exams.

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

1. Comparison between Standard and FIFO Queues:
  - Standard queues offer best-effort ordering, meaning messages could be received out of order in the application.
  - FIFO queues guarantee message ordering, delivering messages exactly in the order they were sent.
  - Standard queues may deliver duplicate messages, while FIFO queues prevent message duplication using explicit deduplication IDs.
2. Performance and Scaling:
  - Standard queues offer nearly unlimited transactions per second, scaling well.
  - FIFO queues provide a default of 300 transactions per second but can achieve up to 3000 messages per second with message batching.
  - FIFO high throughput feature allows processing up to 9000 transactions per second per API call, or up to 90,000 messages per second with batching.
3. Demo: Creating a FIFO Queue:
  - Demonstrated creating a FIFO queue in the AWS Management Console, highlighting configuration options such as content-based deduplication, FIFO throughput limits, and message group IDs.
4. Exam Tips:
  - FIFO queues are ideal for scenarios requiring strict message ordering in distributed decoupled applications.
  - Consider the performance trade-offs between standard and FIFO queues for your application's needs.
  - FIFO queues provide efficient message ordering and deduplication compared to other methods.
  - Understand the use of message deduplication IDs and message group IDs for ensuring ordering and processing limits.
  - Note that FIFO queues may incur higher costs compared to standard queues due to their additional features and benefits.

Overall, understanding FIFO queues is crucial for applications that require strict message ordering and deduplication. Be sure to review these concepts thoroughly, as they are likely to be tested in AWS certification exams

1. Push-Based Messaging with SNS:
  - SNS is a push-based messaging service in AWS, allowing messages sent by a producer to be immediately sent to all subscribed consumers.
  - It enables proactive message delivery to various endpoints such as emails, SQS queues, Lambda functions, and HTTP clients.
2. Settings and Quotas:
  - Messages sent via SNS can be up to 256 kilobytes of text data.
  - SNS supports various subscriber types by default, including SQS queues, Lambda functions, and email/SMS.
  - SNS offers features like encryption at rest using KMS keys, access policies, and leveraging SQS dead-letter queues for failed messages.
3. Large Message Payloads:
  - The SNS Extended Library allows sending messages up to two gigabytes in size by storing the payload in an Amazon S3 bucket and publishing a message containing a reference to it.
4. SNS Fanout:
  - SNS Fanout replicates messages published to a topic to multiple endpoint subscriptions simultaneously, enabling decoupled, parallel, and asynchronous processing.

5. Message Filtering:
  - SNS supports message filtering using JSON policies to define which messages are sent to specific subscribers based on content and attributes.
6. Console Demo:
  - Demonstrated creating an SNS topic with a simple email subscription in the AWS Management Console, confirming the subscription via email, and publishing a message to the topic.
7. Exam Tips:
  - Consider using SNS for real-time alerting and push-based messaging scenarios.
  - Know the supported subscriber types and understand custom retry policies for HTTP/HTTPS endpoints.
  - Understand FIFO SNS topics for message ordering and deduplication.
  - Familiarize yourself with different architectural layouts and scenarios involving SNS.

By understanding these concepts and following the exam tips, you'll be well-prepared for any questions related to Amazon SNS in AWS certification exams or real-world applications.

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

1. What is Amazon API Gateway?:
  - API Gateway is a fully managed service in AWS that enables you to create, publish, and update secured APIs for your applications.
  - It acts as a front door entrance for your applications using RESTful APIs and offers simple integrations with AWS Lambda functions, HTTP endpoints, and other AWS services.
2. Notable Features:
  - Security features include AWS WAF integration, DDoS protection, and request throttling.
  - API Gateway is designed to be easy to use, allowing for customized API creation.
3. API Options:
  - RESTful APIs: Traditional option supporting features like API keys, request validation, and integration with AWS WAF.
  - HTTP APIs: Simplified option with minimal features compared to RESTful APIs, but cost-effective.
  - WebSocket APIs: Support two-way interactive communication sessions for real-time applications like notifications.
4. Endpoint Types:
  - Edge-optimized: Default option leveraging CloudFront for global users.
  - Regional: Suitable for APIs specific to users within a particular region.
  - Private: Accessible only via VPC interface endpoints for special use cases within application architectures.
5. Securing APIs:
  - Implement user authentication at the API level using various methods including IAM roles, Amazon Cognito pools, or custom authorizers with Lambda functions.
  - Implement custom TLS with ACM certificates for edge-optimized and regional endpoints.
  - Leverage AWS WAF for additional DDoS and layer seven protection.
6. API Gateway in Practice:
  - Demonstrated creating a RESTful API using API Gateway to front a Lambda function returning inspirational quotes.
  - Deployed the API and accessed it via a publicly accessible URL.
7. Exam Tips:
  - Understand that APIs serve as front doors for applications and services.
  - Whenever scenarios involve accessing backend services, consider using API Gateway.

- API Gateway supports versioning, WAF integration, and different endpoint types.
- Remember the different API options (RESTful, HTTP, WebSocket) and endpoint types (edge-optimized, regional, private).

By understanding these concepts and following the exam tips, you'll be well-prepared for questions related to Amazon API Gateway in AWS certification exams or real-world application scenarios

#### 1. Overview of AWS Batch:

- AWS Batch is a managed service that simplifies the creation and execution of batched compute-based workloads in AWS.
- It handles scaling based on configured infrastructure and optimizes workload distribution automatically.
- No installation is required, and Batch abstracts the maintenance of batch workload software.

#### 2. Important Components:

- Job: A unit of work submitted to Batch, which can be shell scripts, executable files, or Docker images.
- Job Definition: Specifies how jobs will be run, acting as blueprints for resources within jobs.
- Job Queue: Where jobs are submitted and reside until scheduled to run.
- Compute Environment: Managed or unmanaged compute resources that run jobs.

#### 3. Fargate vs. EC2 Compute Environments:

- Fargate is recommended for workloads requiring fast start times and specific CPU and memory requirements.
- EC2 is suitable for workloads needing more control over instance selection, GPU requirements, high concurrency rates, etc.

#### 4. Comparison of Batch and Lambda:

- Lambda is ideal for short-running workloads (<15 minutes) with limited disk space.
- Batch can run longer tasks and supports Docker, providing more flexibility in runtimes.

#### 5. Architecture Example with Batch:

- Illustrated a workflow where S3 events trigger a Lambda function, which submits jobs to Batch for processing data using EC2, Fargate, or EKS compute environments.

#### 6. Managed and Unmanaged Compute Environments:

- Managed: Batch manages capacity and instance types based on defined specifications, supporting EC2 On-Demand, EC2 Spot, Fargate, and Fargate Spot capacity.
- Unmanaged: Users manage their own compute resources, verifying that the AMI meets Amazon ECS container instance specifications.

#### 7. Exam Tips:

- Understand that AWS Batch is suitable for long-running, event-driven workloads.
- Be familiar with job submission, job definitions, job queues, and compute environments.
- Consider limitations of Lambda (time limit, disk space, runtimes) when deciding between Lambda and Batch.
- Understand the differences between Fargate and EC2 compute environments.
- Know the distinction between managed and unmanaged compute environments, with managed being the preferred choice for most scenarios.

By internalizing these concepts and exam tips, you'll be well-prepared to tackle questions related to executing batch workloads using AWS Batch in certification exams or real-world scenarios

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

#### 1. Overview of Amazon MQ:

- Amazon MQ is a message broker service in AWS designed to facilitate the



migration of existing messaging applications to the cloud.

- It supports multiple programming languages, operating systems, and messaging protocols.
  - The service offers compatibility with Apache ActiveMQ and RabbitMQ engine types.
2. Choosing Between SNS with SQS and Amazon MQ:
    - Both SNS with SQS and Amazon MQ offer architectures with topics and queues for messaging.
    - Amazon MQ is suitable for migrating existing applications that already use messaging systems, while SNS with SQS is preferred for new messaging systems due to its simplicity and scalability.
    - Amazon MQ requires private networking for access within a VPC, Direct Connect, or VPN, while SNS with SQS is publicly accessible by default.
  3. Amazon MQ Brokers:
    - Amazon MQ offers highly available architectures depending on the chosen engine type (ActiveMQ or RabbitMQ).
    - ActiveMQ uses an active-standby architecture with separate maintenance windows, while RabbitMQ employs clustered deployments across availability zones.
  4. Exam Tips:
    - Amazon MQ is ideal for migrating existing message broker systems and supports various messaging protocols like JMS, AMQP, MQTT, and STOMP.
    - Consider using SNS with SQS for new applications due to its simplicity.
    - Ensure private networking is in place to connect to Amazon MQ clusters and deployments.
    - Understand the highly available architectures offered by Amazon MQ based on the selected engine type.

By grasping these concepts and exam tips, you'll be well-prepared to tackle questions related to brokering messages using Amazon MQ in certification exams or real-world scenarios

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

1. Overview of AWS Step Functions:
  - AWS Step Functions is a serverless orchestration service that enables the combination of AWS Lambda functions with other AWS services to build business-critical applications.
  - It offers a graphical console to visualize application workflows and event-driven steps.
2. Components of AWS Step Functions:
  - State Machines: Workflows consisting of event-driven steps, where each step is considered a state.
  - Tasks: States within workflows that represent single units of work performed by AWS services.
3. Types of Workflows:
  - Standard Workflows: Provide exactly-once execution and can run for up to one year, suitable for long-running workflows requiring an auditable history.
  - Express Workflows: Offer at-least-once execution and have a maximum runtime of 5 minutes, ideal for high-event-rate workloads like IoT data ingestion.
4. Integrated AWS Services:
  - AWS Step Functions integrates with various AWS services, including Lambda, SNS, API Gateway, and Fargate, among others.
5. Types of States:
  - Pass State: Passes input directly to output without performing any work.
  - Task State: Represents a single unit of work performed by an AWS service like Lambda or SNS.
  - Choice State: Implements branching logic based on input values.

- Wait State: Introduces a specified time delay within a workflow.
- Succeed State: Stops executions successfully.
- Fail State: Stops executions and marks them as failures.
- Parallel State: Runs parallel branches of executions within state machines.
- Map State: Executes a set of steps based on elements of an input array.

#### 6. Exam Tips:

- Understand the purpose and capabilities of AWS Step Functions, especially its integration with other AWS services.
- Be familiar with the execution types: standard and express workflows.
- Know that workflows are written in Amazon States Language (ASL), a format similar to JSON.
- Recognize the wide range of integrations available and common use cases for Step Functions.
- Understand the different types of states and when to use them, particularly for adding time delays or branching logic.

By mastering these concepts and exam tips, you'll be well-prepared to tackle questions related to coordinating distributed applications with AWS Step Functions in certification exams or real-world scenarios

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>

#### 7. Overview of Amazon AppFlow:

- Amazon AppFlow is a fully managed service that facilitates the secure exchange of data between SaaS applications and AWS services.
- It allows bidirectional data transfer, enabling ingestion of data from SaaS applications into AWS and vice versa.

#### 8. Important Terms and Concepts:

- Flow: The mechanism that transfers data between a source and a destination.
- Data Mapping: Determines how data from the source is mapped to fields in the destination.
- Filters: Control which data records are transferred based on defined criteria.
- Triggers: Specify how flows are initiated, such as on-demand, on event, or on schedule.

#### 9. Example Diagram:

- Illustrated the flow of data from a SaaS application to an S3 bucket in AWS.
- Steps include creating source and destination connections, mapping fields, applying filters, and running the flow.

#### 10. Use Cases:

- Transferring Salesforce records to Amazon Redshift for analysis.
- Ingesting and analyzing Slack conversations and storing them in S3.
- Migrating Zendesk or other help desk support tickets to services like Snowflake.
- Scheduled transfer of aggregate data to Amazon S3, limited to 100 gigabytes per flow.

#### 11. Exam Tips:

- Understand Amazon AppFlow as a managed service for integrating data from SaaS applications into AWS.
- Recognize that AppFlow supports bidirectional data transfer and offers various triggers for flow execution.
- Be aware of use cases where AppFlow can streamline data ingestion processes, especially for scenarios involving scheduled transfers or large amounts of SaaS data.

By grasping these concepts and exam tips, you'll be well-prepared to tackle questions related to ingesting data from SaaS applications into AWS using Amazon AppFlow.

#### 1. Understanding Workloads:

- Distinguish between synchronous and asynchronous workloads to determine suitable AWS services.
- 2. Decoupling Strategies:
  - Consider the type of decoupling needed, such as Pub/Sub with SNS/SQS or workflow orchestration with Step Functions.
- 3. Message Ordering and Duplication:
  - Assess whether message order matters and be aware of message duplication possibilities in services like SQS.
- 4. Service Defaults and Limits:
  - Understand default settings and service limits for AWS services like SQS and Step Functions.
- 5. Service-Specific Tips:
  - Amazon SQS: Know about message duplication, queue directionality, defaults, and ordering options.
  - Amazon SNS and API Gateway: Use SNS for push-based notifications, integrate with CloudWatch alarms, and understand API Gateway's role as a secure front door.
  - AWS Batch: Suitable for long-running batch workloads and provides an alternative to AWS Lambda for certain use cases.
  - Amazon MQ: Managed messaging broker service supporting RabbitMQ and ActiveMQ engines, suitable for specific messaging protocols like JMS.
  - AWS Step Functions: Serverless orchestration service for workflow automation, ideal for scenarios requiring decision logic and lengthy wait periods.
  - Amazon AppFlow: Bi-directional service for ingesting data from third-party SaaS applications into AWS services, suitable for regular data ingestion scenarios.

Remember these tips when approaching exam questions related to decoupling workflows and AWS services. If you understand the characteristics, use cases, and limitations of each service, you'll be better prepared to choose the most appropriate solution for different scenarios.

From <<https://chat.openai.com/c/d6fe90a1-a00b-4374-9bfc-54c5891880f3>>