


- Data processing



LEARNING OBJECTIVES

At the end of this unit, you should be able to:

- Batch Processing
- Real Time Processing
- Batch Vs Real time processing



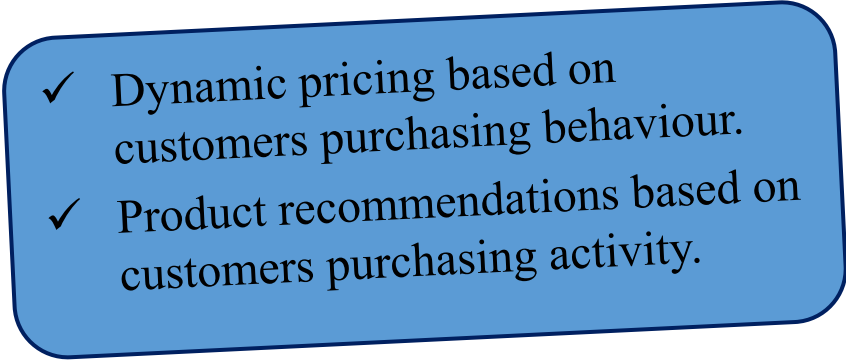
EVERY NOW AND THEN, ALMOST IN EVERY CORNER OF THE WORLD..... THIS HAPPENS...!!



EVERY NOW AND THEN, ALMOST IN EVERY CORNER OF THE WORLD..... THIS HAPPENS...!!

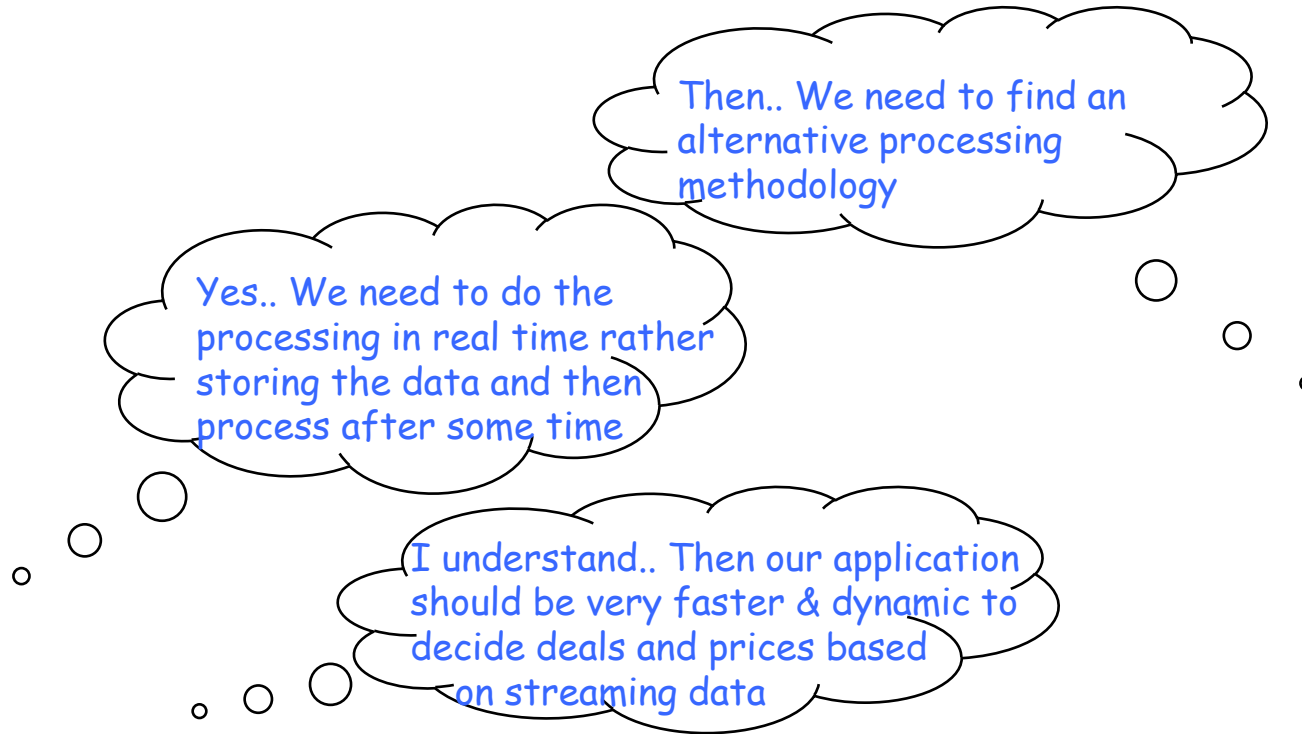


Team...! We need to implement the following features in upcoming Great Indian Sale.

- 
- ✓ Dynamic pricing based on customers purchasing behaviour.
 - ✓ Product recommendations based on customers purchasing activity.

AT THE SAME TIME.... IN AN ECOMMERCE APPLICATION DEVELOPMENT & ANALYTICS TEAM

•




LOOKING AT THE SCENARIO, WHAT IS REQUIRED?

Read More in Database Storage & Design View All & Manage
Upcoming Deals exclusive for you in Database Storage & Design



Big Data Analytics
Parag Kulkarni
★★★★★ 1
₹ [] prime



Machine Learning for Big Data:
Hands-On for Developers and
Technical...
Jason Bell
★★★★★ 9
₹ [] prime



Introducing Data Science
Data, Machine Learning,
More...
Davy Clalen
★★★★★ 3
₹ [] prime

From Shopping website

1

Login ID : C003, Location : Bangalore
Search Page: Data Analytics books
Time: 16:23:45.34, device : mobile
IP: 10.0.12.13

Update price in webpages

Big Data Analytics : 566 (15 % off)
Machine Learning for Big Data : 449 (43% off)
Introduction to data science : 687 (12% off)

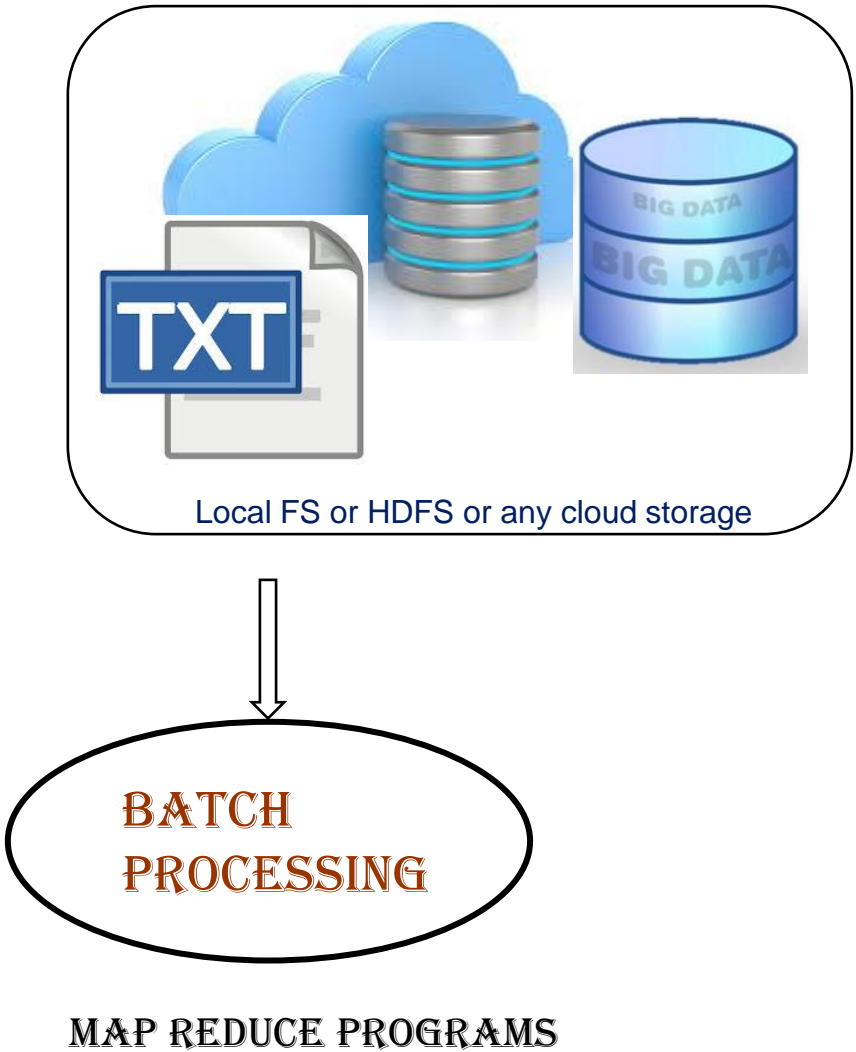
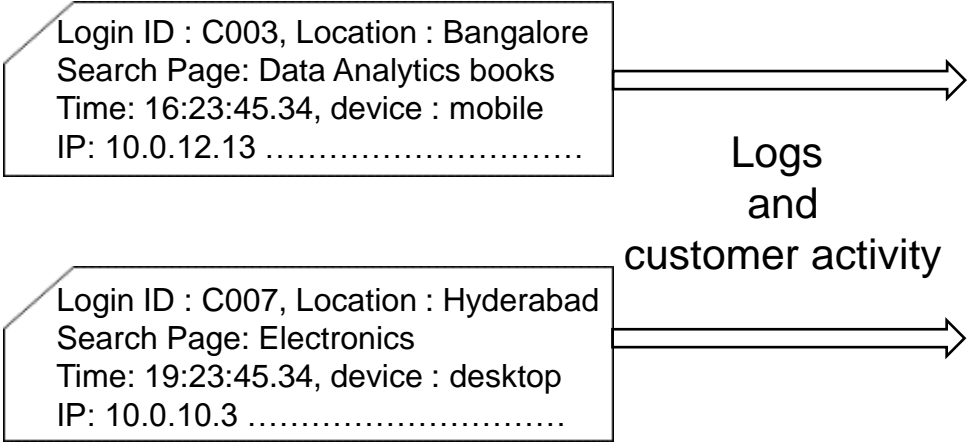
3

Application to decide the
price of books based on
customer logs and
previous search activities

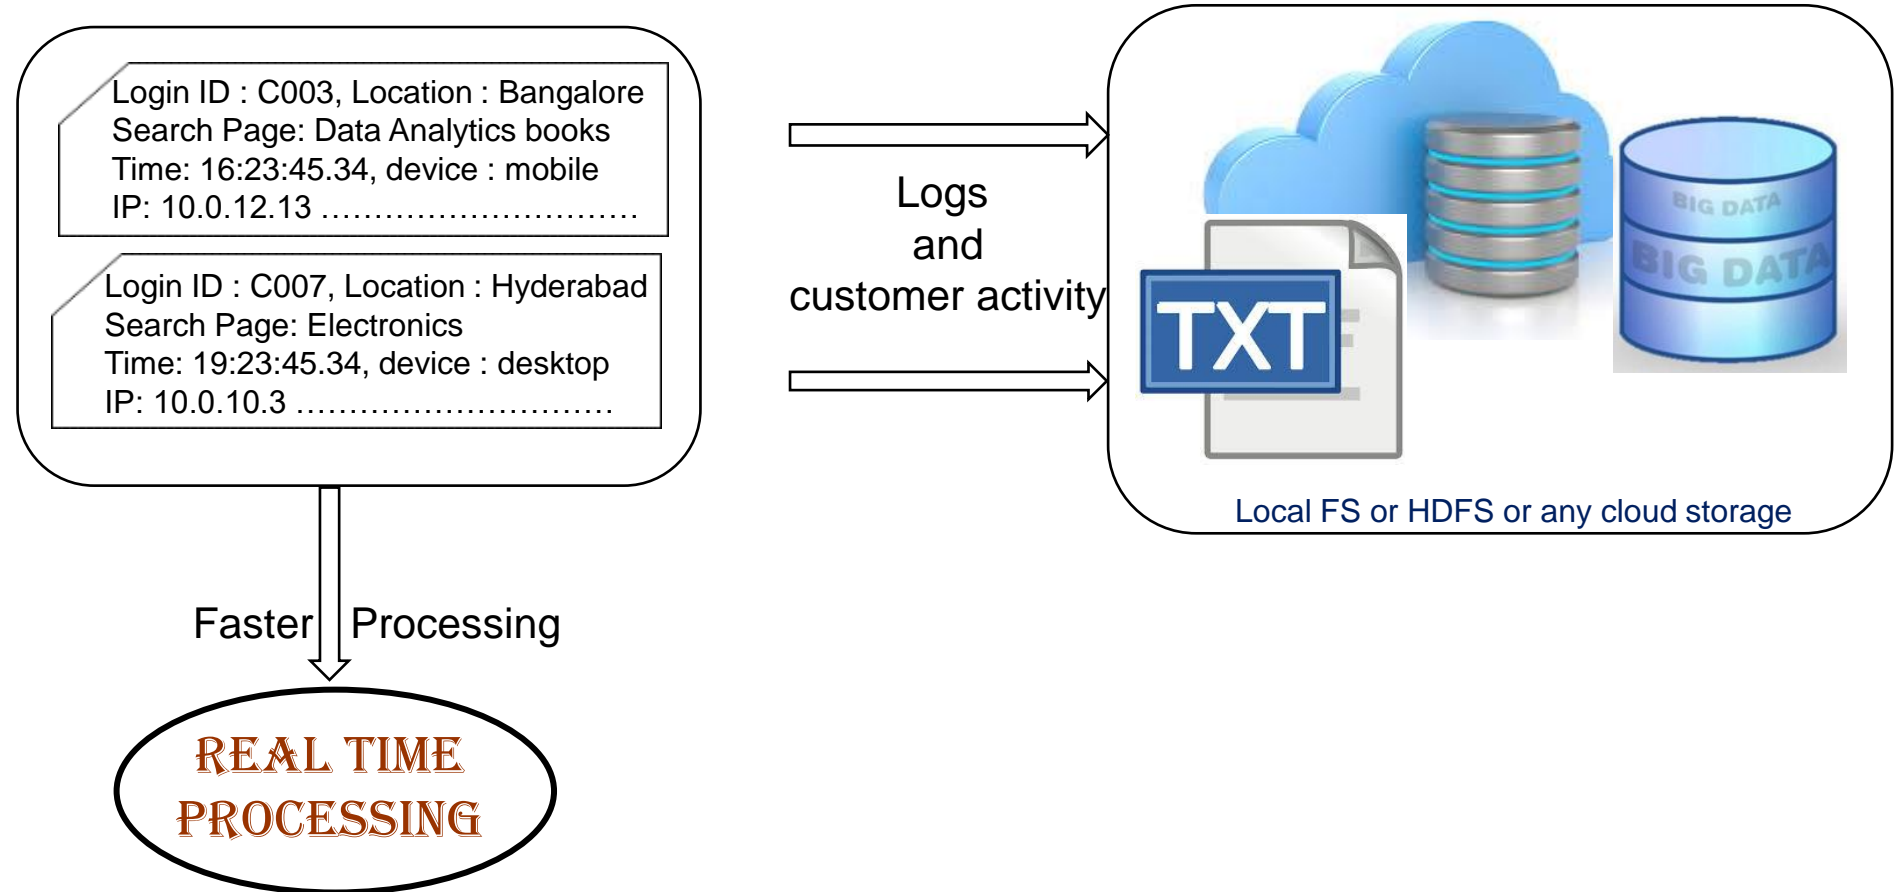
2

Previous activity logs of customer
C003 stored in HDFS or
Local FS or any other storage

WHAT WE HAVE DONE SO FAR ?



WHAT WE NEED TO DO?




HERE WE HAVE...

APACHE SPARK

The Framework for
Faster Batch & Real Time Data Processing

- Apache Spark



LEARNING OBJECTIVES

At the end of this unit, you should be able to:

- What is Spark?
- Why Spark?
- Spark Vs Hadoop



What is Spark ?

- Open source fast and general engine for large data processing.
- Developed in 2009 in UC Berkeley's AMP Lab, and open sourced in 2010 as an Apache project.
- Written in Scala.
- In memory processing framework.
- Provides high-level APIs in Java, Scala, Python and R.

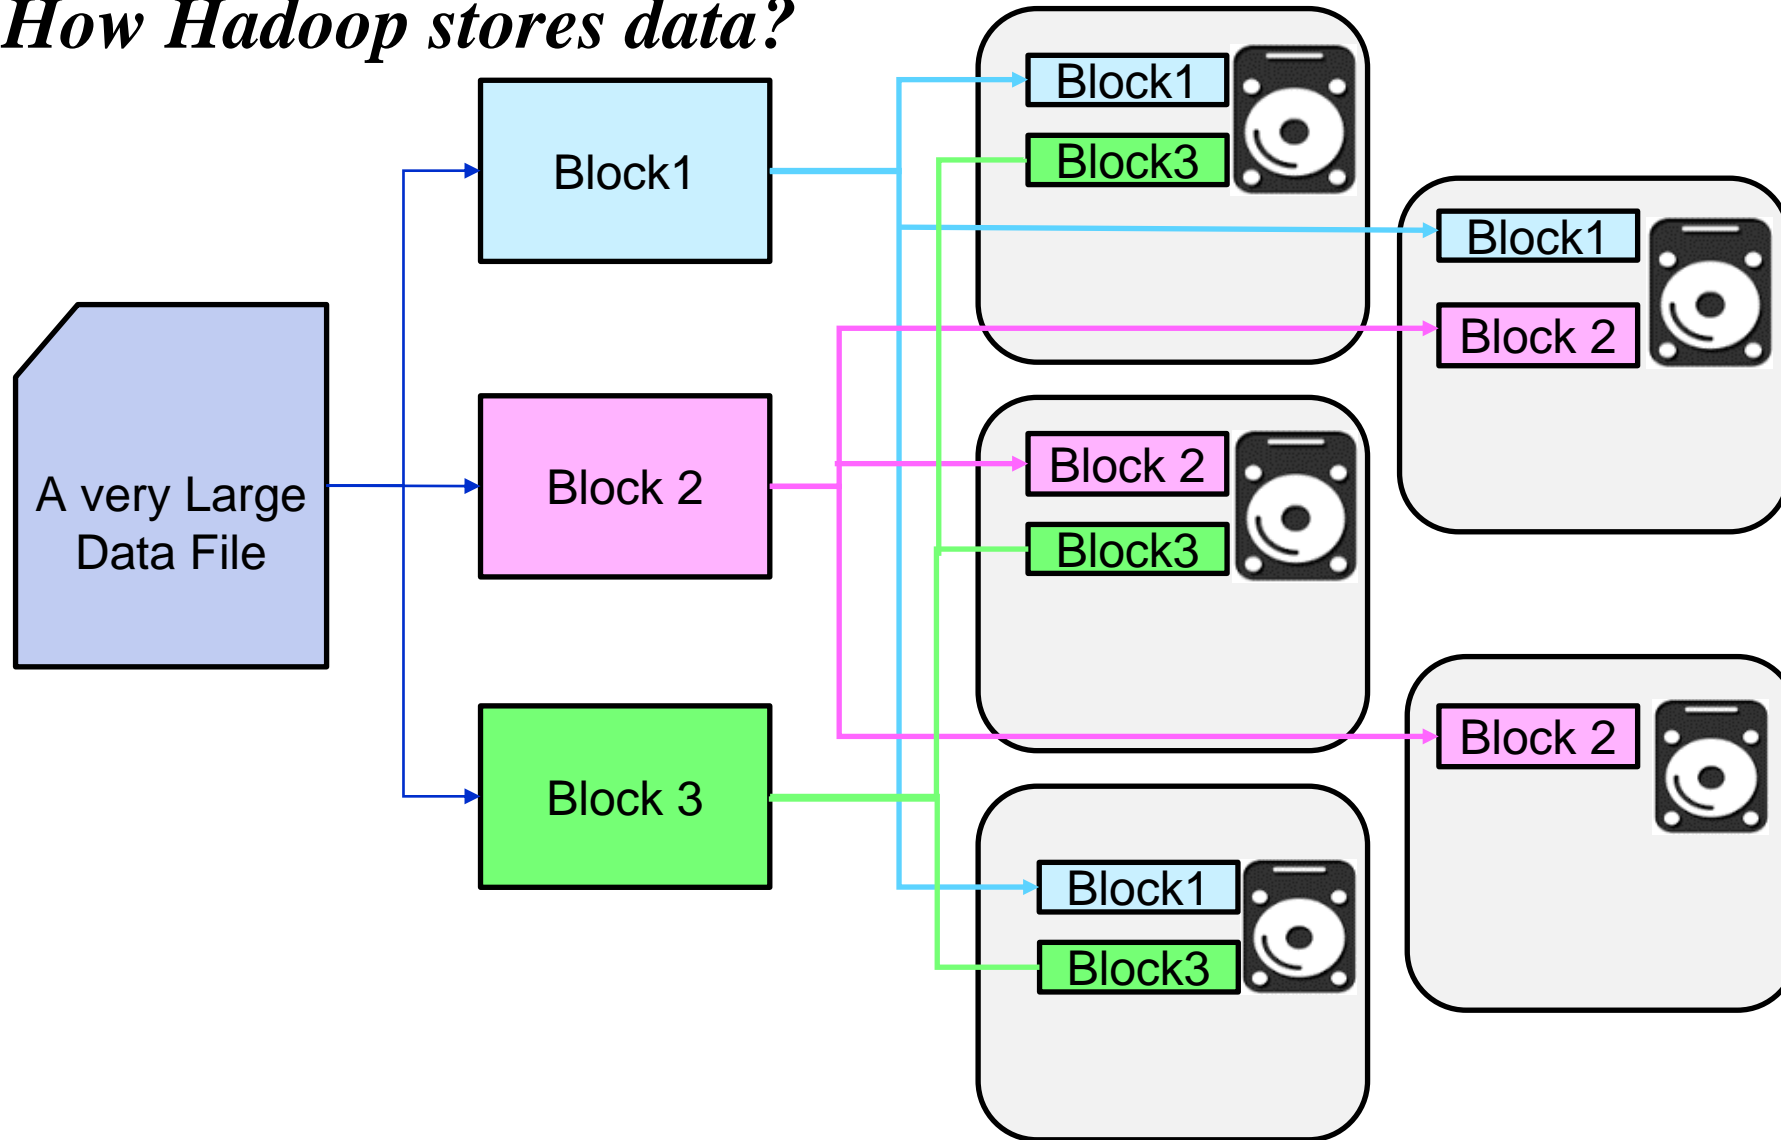


Why Spark ?

- Unified framework to manage big data requirements.
- Provides high level operators such as filter, map, etc..
- 100x faster in memory, 10x faster on disk.
- Spark Shell
 - - Interactive - for data exploration and testing.
 - - Scala or Python
- Spark Applications
 - - For large scale and data processing.
 - - Scala, Python or Java

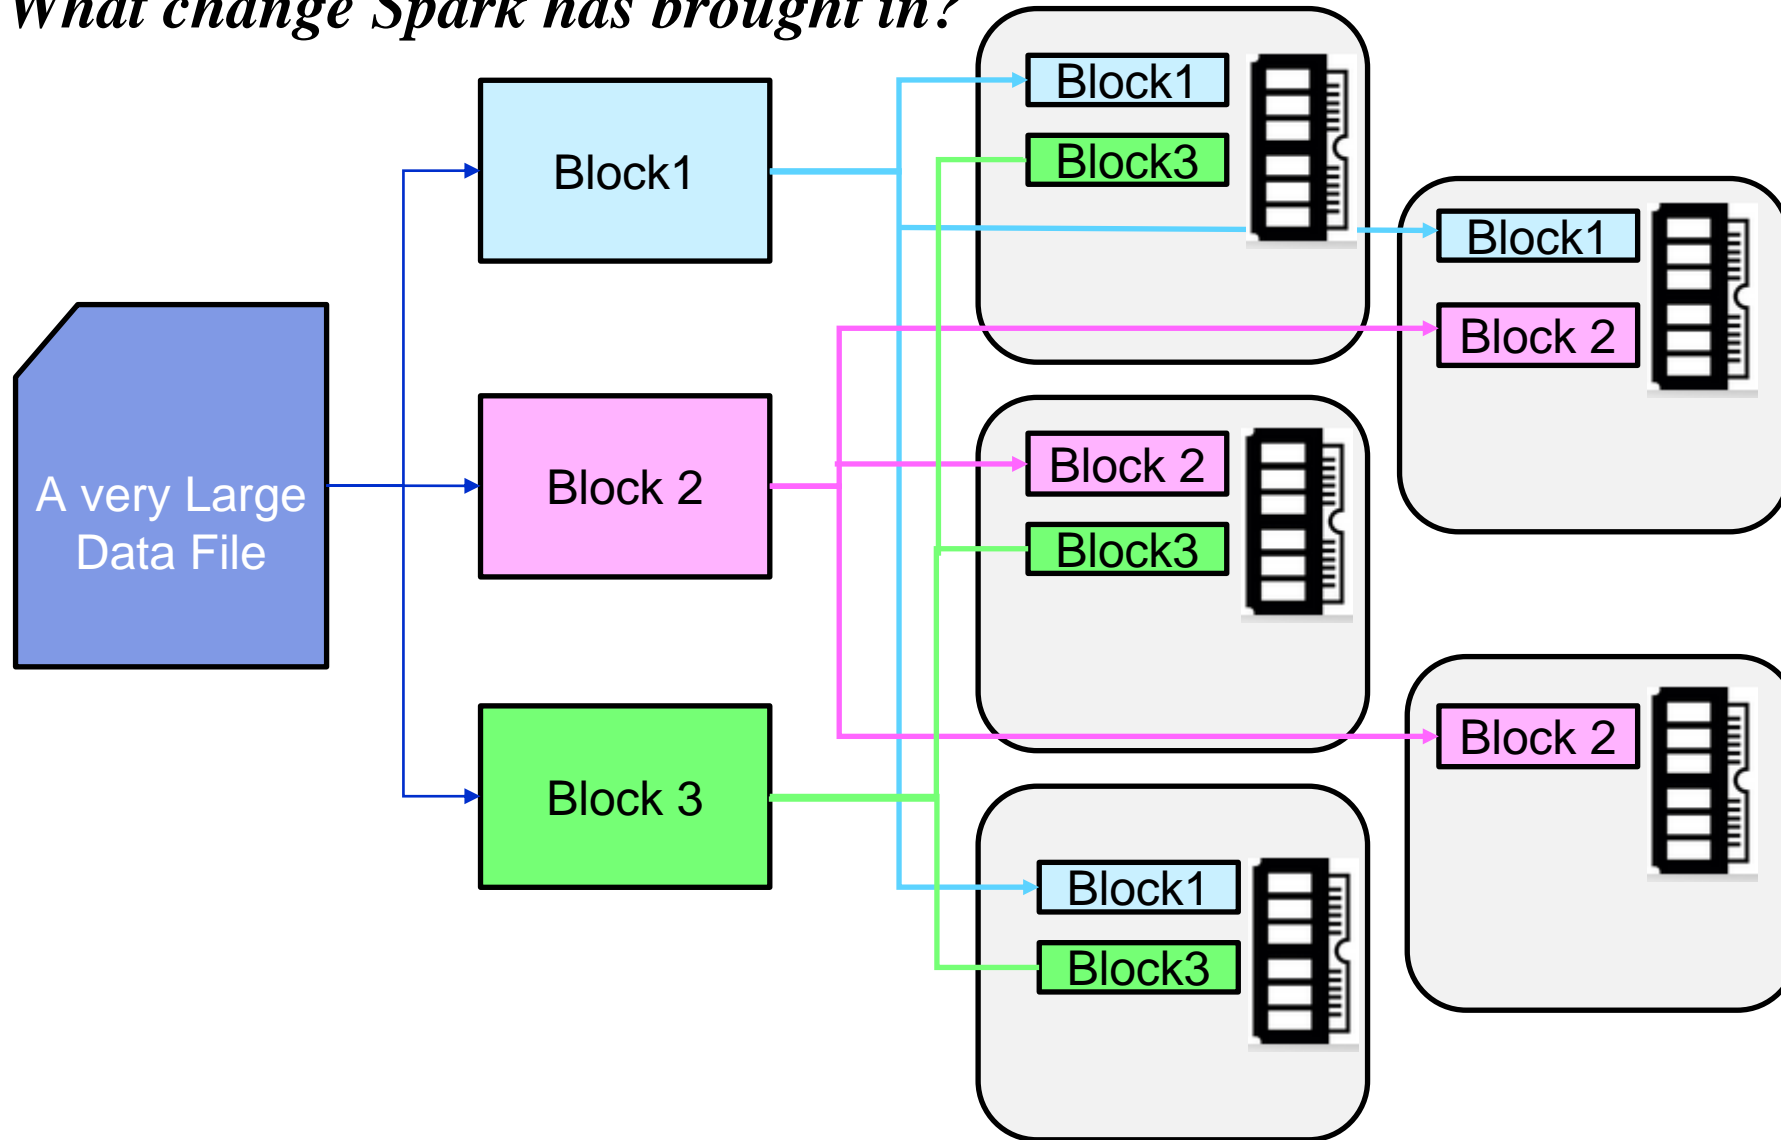
Spark Vs Hadoop

How Hadoop stores data?



Spark Vs Hadoop

What change Spark has brought in?



Spark Vs Hadoop

	Spark	Hadoop
Introduction	<ul style="list-style-type: none">▪ Faster and general purpose data processing engine.▪ Handles batch as well real time processing.	<ul style="list-style-type: none">▪ Processes structured and unstructured data that are stored in HDFS.▪ Handles only batch processing.

Spark Vs Hadoop

	Spark	Hadoop
Speed	<ul style="list-style-type: none">▪ Run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk.▪ Reducing the number of read/write cycle to disk and storing intermediate data in-memory.	<ul style="list-style-type: none">▪ Reads and writes from disk and that slows down the processing speed.

Spark Vs Hadoop

	Spark	Hadoop
Difficulty	<ul style="list-style-type: none">▪ Provides high level operators. E.g. filter, map	<ul style="list-style-type: none">▪ No high level operators.▪ Need to hand code each and every operation.

Spark Vs Hadoop

Big Code




```
public class WordCount {  
    public static class TokenizerMapper  
        extends Mapper<Object, Text, Text, IntWritable>{  
  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(Object key, Text value, Context context  
            ) throws IOException, InterruptedException {  
            StringTokenizer itr = new StringTokenizer(value.toString());  
            while (itr.hasMoreTokens()) {  
                word.set(itr.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
  
    public static class IntSumReducer  
        extends Reducer<Text, IntWritable, Text, IntWritable> {  
        private IntWritable result = new IntWritable();  
  
        public void reduce(Text key, Iterable<IntWritable> values,  
            Context context  
            ) throws IOException, InterruptedException {  
  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            result.set(sum);  
            context.write(key, result);  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "word count");  
        job.setJarByClass(WordCount.class);  
        job.setMapperClass(TokenizerMapper.class);  
        job.setCombinerClass(IntSumReducer.class);  
        job.setReducerClass(IntSumReducer.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Tiny Code



```
val textFile = sc.textFile("hdfs://...")  
val counts = textFile.flatMap(line => line.split(" "))  
                        .map(word => (word, 1))  
                        .reduceByKey(_ + _)  
counts.saveAsTextFile("hdfs://...")
```

- Spark Architecture



LEARNING OBJECTIVES

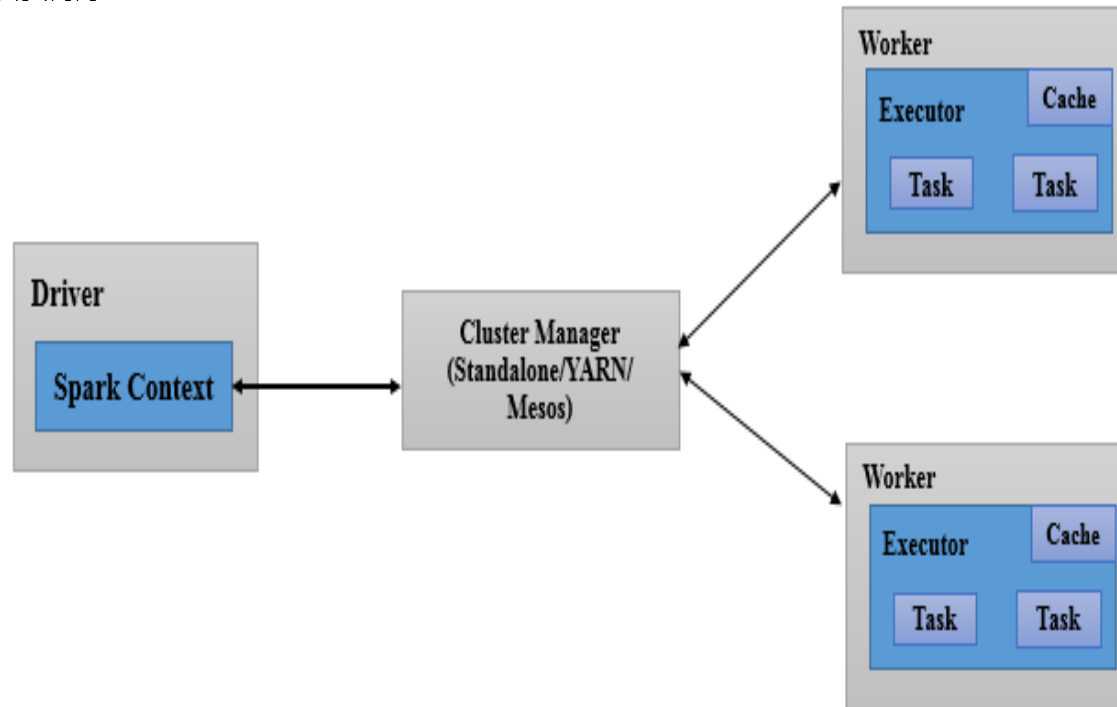
At the end of this unit, you should be able to:

- Architecture Overview
- Spark Cluster setup – Pseudo distributed mode
- Spark Shell
- Spark Context



ARCHITECTURE OVERVIEW

- Master/Slave architecture with cluster manager and two daemons.
- Daemons are:
 - Master – Master/ Driver Process
 - Worker – Slave Process



ARCHITECTURE OVERVIEW

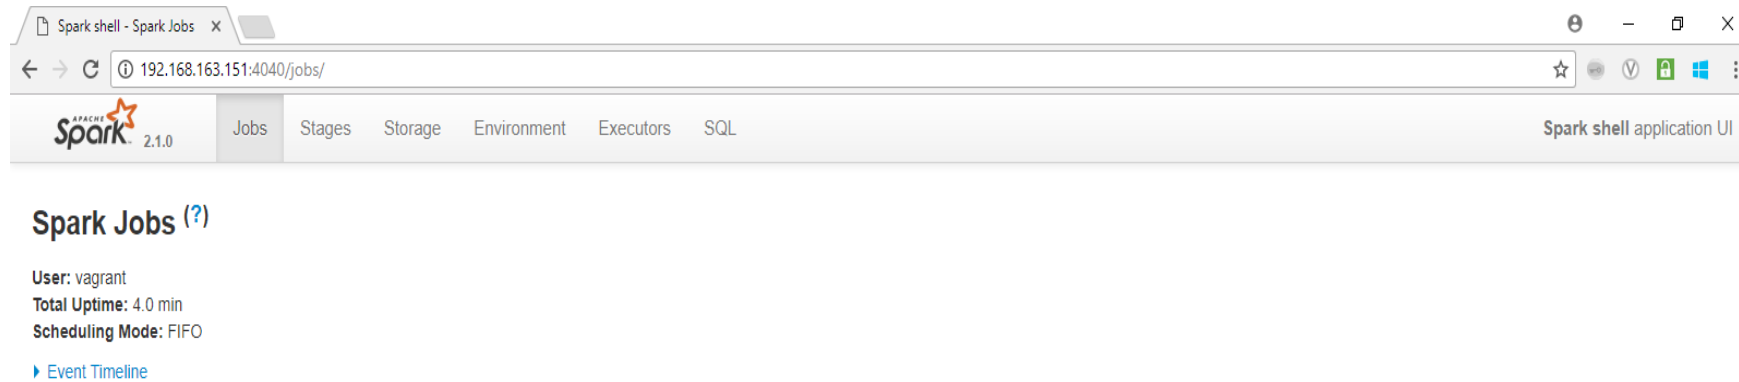
- A spark cluster has single coordinator called driver and many distributed workers.
- Driver communicate with large number of distributed workers called executors.
- Cluster Manager is responsible for scheduling and allocating resources to a Spark Job.

ARCHITECTURE OVERVIEW - ROLE OF DRIVER

- The driver program
 - Entry point of the Spark Shell (Scala, Python, and R).
 - Runs application main () function
 - Is the place where Spark Context is created.
- Responsible for:
 - Scheduling job, negotiating with the cluster manager.
 - Converting a user application into smaller execution units i.e. tasks.
- You can access running spark application information through default Web UI at port 4040.

ARCHITECTURE OVERVIEW - ROLE OF DRIVER

- **Spark shell application UI.**



ARCHITECTURE OVERVIEW - ROLE OF EXECUTORS

- Distributed agent.
- Responsible for:
 - Executing tasks.
 - Performing data processing.
 - Reading from and Writing data to external sources.
 - Storing computation results data in-memory, cache or disk.
 - Interacting with the storage systems.

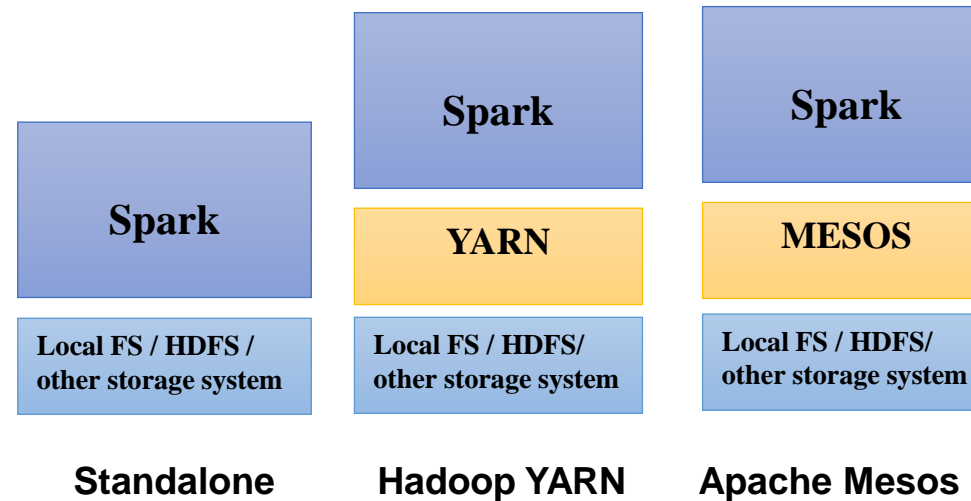
ARCHITECTURE OVERVIEW - ROLE OF CLUSTER MANAGER

- It is an external service.
- Responsible for:
 - acquiring resources and allocating them to a spark application.
 - allocation and deallocation of various physical resources such as CPU, memory, etc.,

ARCHITECTURE OVERVIEW - ROLE OF CLUSTER MANAGER

.Three types of cluster managers:

- Standalone Cluster Manager
- Hadoop YARN
- Apache Mesos



SPARK CLUSTER SETUP

Pseudo distributed mode cluster setup

SPARK SHELL

- Provides interactive shell for data exploration and testing (REPL).
- To start Spark Shell, type `spark-shell` command as shown below.

Note: REPL (Read/Evaluate/Print Loop)

```
vagrant@master:~$ spark-shell
```

```
Spark context Web UI available at http://192.168.163.151:4040
Spark context available as 'sc' (master = local[*], app id = local-1509877354288).
Spark session available as 'spark'.
```

Welcome to

[illegible]


```
Using Scala version 2.11.8 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_77)
Type in expressions to have them evaluated.
Type :help for more information.
```

```
scala>
```

SPARK CONTEXT


- At high level, every Spark application requires Spark Context.
- Spark Context is entry point to the Spark API.
- Driver program communicates with Spark through Spark Context.

```
Spark context available as 'sc' (master = local[*], app id = local-1509877354288).  
Spark session available as 'spark'.  
Welcome to
```



```
scala> sc.appName  
res0: String = Spark shell  
  
scala> █
```

- Spark Eco System



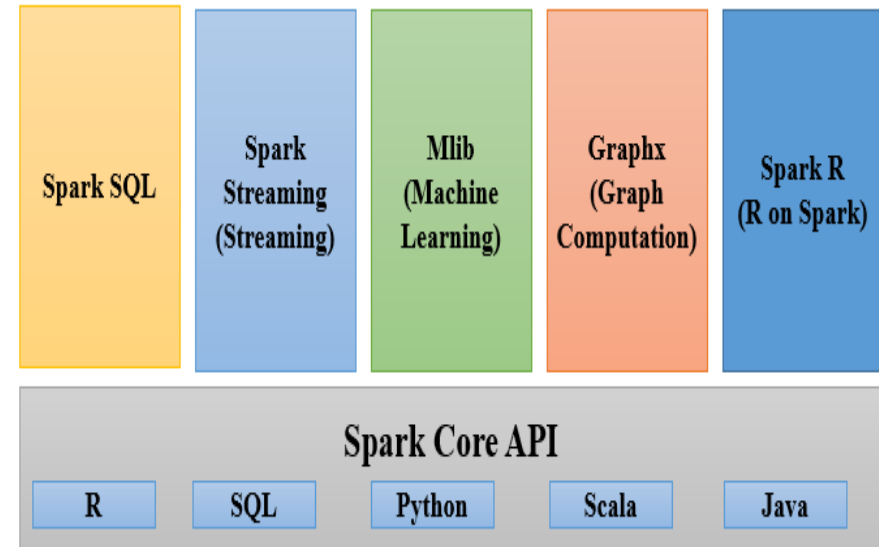
Spark Eco System

Spark Core(RDD):

- Fundamental data structure of Spark.
- RDD (Resilient Distributed Dataset), fault-tolerant collection of elements that can be operated on in parallel.

Spark SQL:

- Running SQL like queries on Spark data.



Spark Eco System

Spark Streaming:

- Scalable, fault-tolerant, high stream processing of live data streams.

MLib:

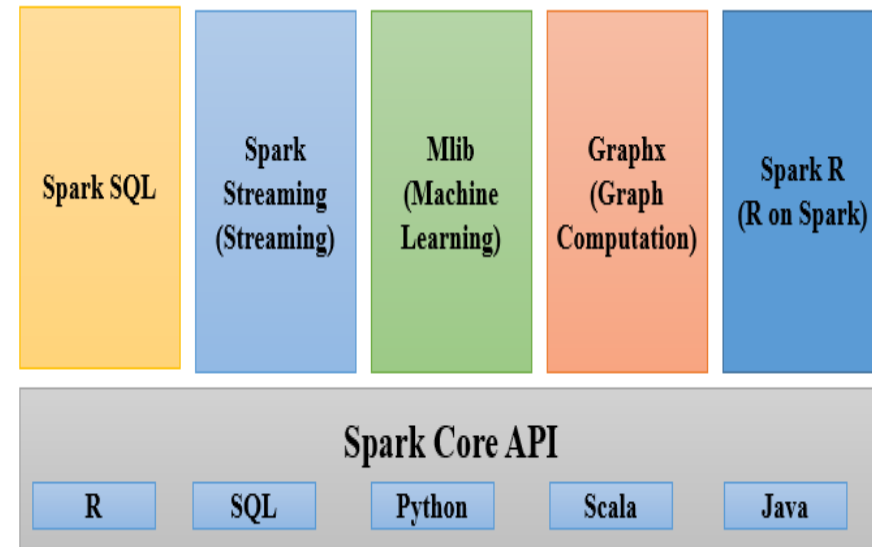
- Scalable practical machine learning.

Graphx:


- Graphs and graph-parallel computation.

SparkR:

- R package that provides a light-weight frontend to use Apache Spark from R.



- Introduction to RDD



LEARNING OBJECTIVES

At the end of this unit, you should be able to:

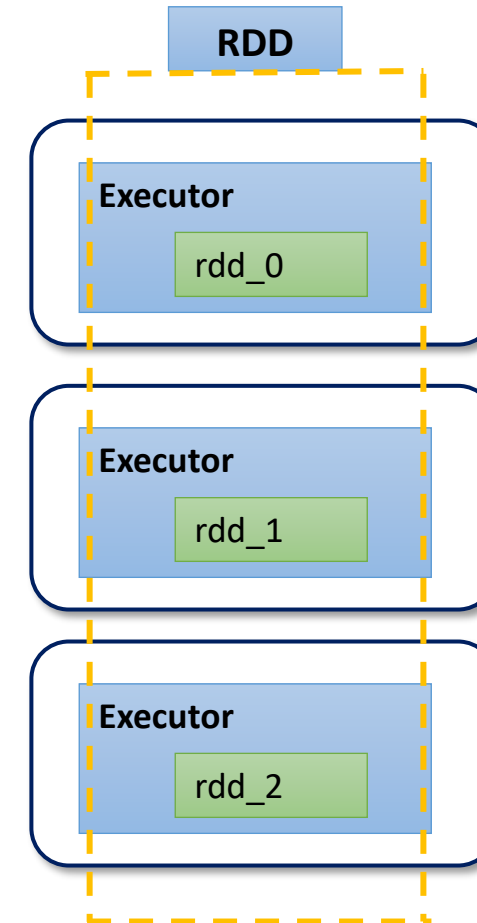
- What is RDD?
- RDD Characteristics
- Partitions
- Read - Only
- RDD operations
- Creating an RDD



WHAT IS RDD?

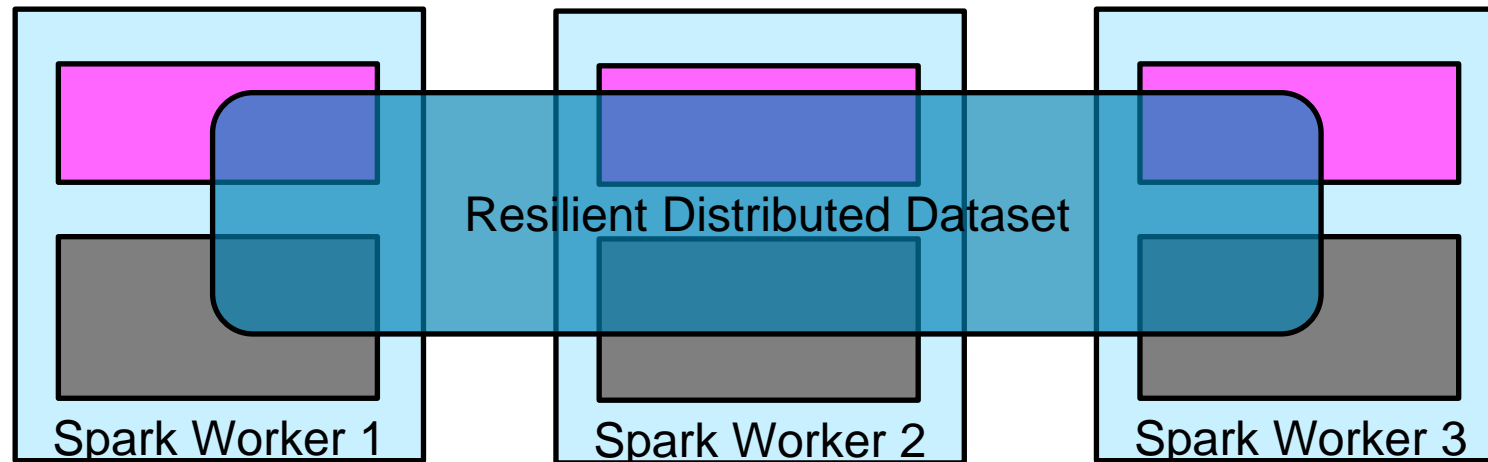
RDD (Resilient distributed dataset)

- Fundamental data structure of Apache Spark.
- Resilient - In built fault tolerance. If something goes wrong reconstruct from source (Lineage).
- Distributed - data is distributed in memory across the worker nodes.
- Dataset - represents records of the data.



WHAT IS RDD?

- RDD is an immutable collection of objects which computes on the different node of the cluster.
- Dataset in RDDs are logically partitioned across many nodes for parallel computation.



RDD CHARACTERISTICS

Partitions:

- RDD represent data in memory.
- To handle huge volume of data
 - data is divided in to partitions that are distributed to multiple machines called nodes.
 - process data in parallel.

RDD CHARACTERISTICS

Ready - Only:

- RDDs are immutable.
 - can not be modified in place.
 - to modify, there are only two types of operations.
 - Transformations
 - Actions


RDD OPERATIONS

Transformation:

- Converts an RDD into another RDD.
- Transform records in dataset.
- Once a dataset is loaded into memory, you can perform chain of transformation before you see some results.

Example:

- Filtering out only certain records.
- Extracting specific field.



RDD OPERATIONS

Actions:

- The actual data processing doesn't happen immediately.
- It happens only user requests a result.

Example:

- First 10 rows
- A Sum
- A Count



CREATING AN RDD

- Two ways to create an RDD.
 - by loading an external dataset, or
 - by distributing a collection of objects (e.g., a list or set).

CREATING AN RDD - USING PARALLELIZED COLLECTIONS

Parallelized Collections:

Objective: Creating an RDD using parallelized collection method of Spark Context.

Action:

```
val input = Array(1, 2, 3, 4, 5)
val output = sc.parallelize(input)
output.collect() // To display output on the console
```

Note: collect() is an Action

Output:

```
scala> val input = Array(1, 2, 3, 4, 5)
input: Array[Int] = Array(1, 2, 3, 4, 5)

scala> val output = sc.parallelize(input)
output: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:26

scala> output.collect()
res1: Array[Int] = Array(1, 2, 3, 4, 5)
```

CREATING AN RDD – USING TEXT FILE

File based RDD:

TextFile RDD can be created using textFile method of Spark Context.

Input Data:

File Name: "goShopping_WebClicks.dat " (contains user activity)

12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=newarrivals	283	google.co.jp	android
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=topbrands	19	diigo.com	windows
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=allwesternwear	18	accuweather.com	android
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=shirts	25	cloudflare.com	windows
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=tops_tees	36	harvard.edu	android
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=dresses	57	simplemachines.org	windows
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=jeans_trousers	62	edublogs.org	mac
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=allethnicwear	22	prnewswire.com	mac
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=kurtas	16	chronoengine.com	windows
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=salwarsuits	9	pinterest.com	linux
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=sarees	37	epa.gov	linux
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=lingerie	107	youku.com	windows
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=sleep_lounge	125	github.com	mac
12/01/2016	17:16:13	128.230.247.37	get	215.82.23.2	https://www.goshopping.com/?product=clothing&producttype=sportswear	37	nature.com	mac

CREATING AN RDD – USING TEXT FILE

Input Data:

File Name: "goShopping_IpLookup.txt" (contains user IP details)

```
172.189.252.8,UK,J5,Devon,38.955855,-77.447819
215.82.23.2,UK,J6,Dorset,39.961176,-82.998794
98.29.25.44,UK,J7,Down,41.49932,-81.694361
68.199.40.156,UK,J8,Dumfries And Galloway,40.657602,-73.583184
155.100.169.152,UK,J9,Dunbartonshire,40.760779,-111.891047
38.68.15.223,UK,J5,Devon,32.776664,-96.796988
70.209.14.54,UK,J6,Dorset,27.950575,-82.457178
74.111.6.173,UK,J7,Down,38.87997,-77.10677
128.230.122.180,USA,NM,NewMexico,43.048122,-76.147424
128.122.140.238,USA,NV,Nevada,40.712784,-74.005941
56.216.127.219,USA,NY,NewYork,35.77959,-78.638179
54.114.107.209,USA,NJ,NewJersey,40.728157,-74.077642
74.111.18.59,USA,NM,NewMexico,43.048122,-76.147424
```


CREATING AN RDD – FILE PARTITIONING

Partitions - single file

- Partitions is done based on size.
- You can specify the number of partitions as shown below.

`textFile(filename, minPartitions)`

- The default number of partitions: 2
- more number of partitions = more parallelization



```
sc.textFile("goShopping_WebClicks.txt",1)
```

CREATING AN RDD – USING TEXT FILE

Objective: Create an RDD for user activity (Use Local File System).

Action:

```
val LocalFSRdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
// URL of the local FS  
LocalFSRdd.count() // count() is an action, counts the number of records.
```

Output:

```
scala> val LocalFSRdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
LocalFSRdd: org.apache.spark.rdd.RDD[String] = /home/vagrant/dataset/goShopping_WebClicks.dat MapPartitionsRDD[1] at textFile at <console>:24  
  
scala> LocalFSRdd.count()  
res0: Long = 71492
```


CREATING AN RDD – USING TEXT FILE

Objective: Create an RDD for user activity (Use HDFS File System).

Action:

```
val HDFSRRdd =  
sc.textFile("hdfs://192.168.163.151:9000/data/goShopping_WebClicks.dat") // URL  
of the HDFS path.  
HDFSRRdd.count()
```

Output:

```
scala> val HDFSRRdd = sc.textFile("hdfs://192.168.163.151:9000/data/goShopping_WebClicks.dat")  
HDFSRRdd: org.apache.spark.rdd.RDD[String] = hdfs://192.168.163.151:9000/data/goShopping_WebClicks.d  
at MapPartitionsRDD[3] at textFile at <console>:24  
  
scala> HDFSRRdd.count()  
res1: Long = 71492
```

CREATING AN RDD – USING TEXT FILE

Objective: Create an RDD for user details and activity.

Action:


```
val FilesRdd = sc.textFile("/home/vagrant/dataset")  
FilesRdd.count() // count () is an action.
```

Note: Directory dataset contains user activity and user details.

Output:

```
scala> val FilesRdd = sc.textFile("/home/vagrant/dataset")  
FilesRdd: org.apache.spark.rdd.RDD[String] = /home/vagrant/dataset MapPartitionsRDD[7] at textFile at <console>:24  
  
scala> FilesRdd.count()  
res3: Long = 71515
```

- Transformations and Actions



LEARNING OBJECTIVES

At the end of this unit, you should be able to:

- Transformations
- Actions
- What is Pair RDDs
- Creating Pair RDDs
- Persisting RDDs
- Storage Level
- Accumulators
- Broadcast Variables



Actions

- Actions return values to the driver program.
- Common actions are:
 - `reduce(func)`
 - `collect()`
 - `count()`
 - `first()`
 - `take(n)`
 - `foreach(println)`
 - `saveAsTextFile(path)`



Actions – Collect()

Objective: Display the records of goShopping_WebClicks.dat file.

Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")
rdd.collect()
```

Output:

```
scala> rdd.collect()
res4: Array[String] = Array("12/01/2016 17:16:13      128.230.247.37  get    215.82.23.2    https://www.goshopping.com
/?product=clothing&producttype=newarrivals      283    google.co.jp    android ", "12/01/2016 17:16:13      128.230.24
7.37  get    215.82.23.2    https://www.goshopping.com/?product=clothing&producttype=topbrands      19    diigo.comw
indows ", "12/01/2016 17:16:13      128.230.247.37  get    215.82.23.2    https://www.goshopping.com/?product=clothing&producttype=allwesternwear      18    accuweather.com android ", "12/01/2016 17:16:13      128.230.247.37  get    21
5.82.23.2    https://www.goshopping.com/?product=clothing&producttype=shirts 25    cloudflare.com windows ", "12/01/
2016 17:16:13      128.230.247.37  get    215.82.23.2    https://www.goshopping.com/?product=clothing&producttype=t
ops tees      36    harvard.edu    android ", "12/01/2016 17:16:13      128.230.247....
```

Actions – Count()

Objective: Count the number of records of goShopping_WebClicks.dat file.

Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
rdd.count()
```

Output:

```
scala> val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
rdd: org.apache.spark.rdd.RDD[String] = /home/vagrant/dataset/goShopping_WebClicks.dat MapPartitionsRDD[13] at textFile at  
<console>:24  
  
scala> rdd.count()  
res6: Long = 71492
```

Actions – take(n)

Objective: Display first 2 records of goShopping_WebClicks.dat file.

Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
rdd.take(2)
```

Output:

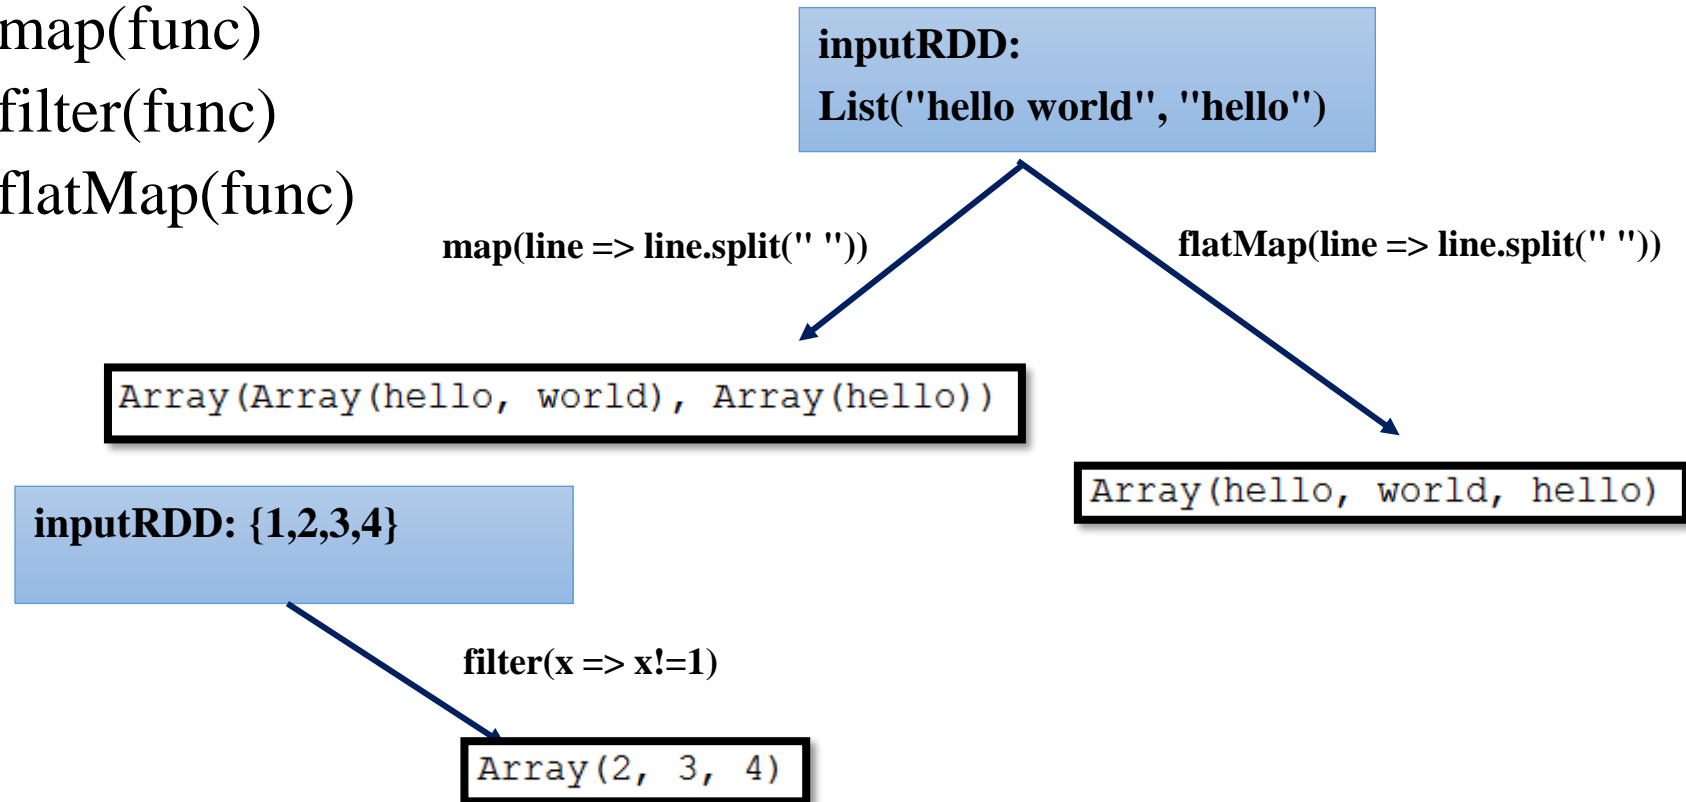
```
scala> rdd.take(2)  
res0: Array[String] = Array("12/01/2016 17:16:13      128.230.247.37  get      215.82.23.2      https://  
/www.goshopping.com/?product=clothing&producttype=newarrivals  283      google.co.jp      android ", "12/  
01/2016 17:16:13      128.230.247.37  get      215.82.23.2      https://www.goshopping.com/?product=clo  
thing&producttype=topbrands      19      diigo.com      windows ")
```


Transformations

- Creating a new RDD from an existing RDD.

- Common transformations are:

- `map(func)`
- `filter(func)`
- `flatMap(func)`



Transformations – map(func)

Objective: Retrieve date, customer ip, time spent on shopping site .

Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
val split = rdd.map(line => line.split("\t"))  
val result = split.map(field => (field(0),field(4),field(6)))  
result.take(5).foreach(println)
```

Output:

```
(12/01/2016,215.82.23.2,283)  
(12/01/2016,215.82.23.2,19)  
(12/01/2016,215.82.23.2,18)  
(12/01/2016,215.82.23.2,25)  
(12/01/2016,215.82.23.2,36)
```

Transformations – filter(func)

Objective: Retrieve date, customer ip, time spent on shopping site only on "12/01/2016" .

Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")
val filtered = rdd.filter(line => line.split("\t")(0).equals(("12/02/2016")))
val split = filtered.map(line => line.split("\t"))
val result = split.map(field => (field(0),field(4),field(6)))
result.take(5).foreach(println)
```

Output:

```
(12/02/2016,155.100.169.152,6)
(12/02/2016,155.100.169.152,6)
(12/02/2016,155.100.169.152,10)
(12/02/2016,155.100.169.152,10)
(12/02/2016,155.100.169.152,10)
```

Lazy Evaluation

- All transformations are lazy, in that they do not compute their results until an action is called.

```
> val rdd = sc.textFile("sample.txt")  
  
    val rdd_uc = rdd.map(line =>  
        line.toUpperCase())  
  
    rdd_uc.collect()
```

Sample.txt

Welcome to spark course.
Spark and Scala.
Spark Programming.
rdd



RDD: rdd

Welcome to spark course.
Spark and Scala.
Spark Programming.
rdd.



RDD: rdd_uc

WELCOME TO SPARK COURSE.
SPARK AND SCALA.
SPARK PROGRAMMING
RDD

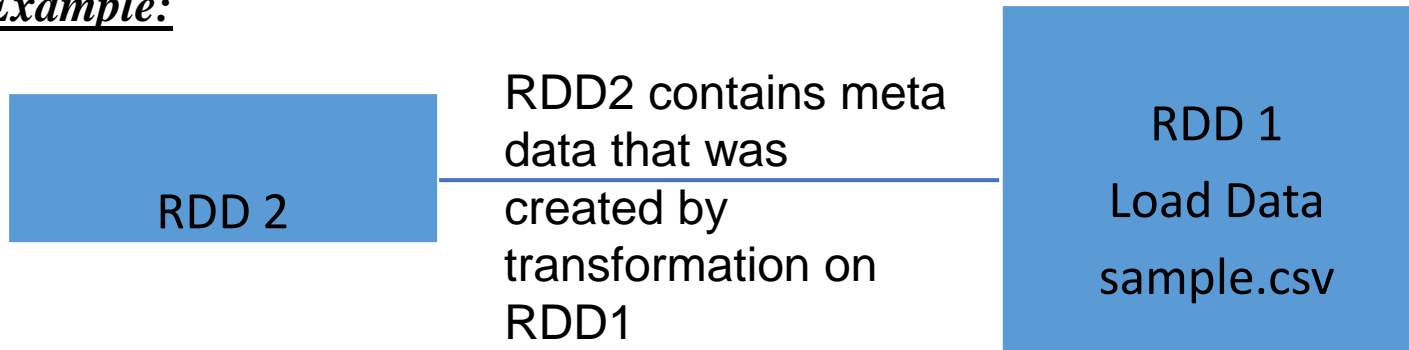


Lazy Evaluation

RDD Lineage and toDebugString():

- Spark maintains each RDDs lineage i.e. previous RDD which it depends.
- When an RDD is created, it just holds metadata.
 - A transformation that created the RDD.
 - Its parent RDD from which it was created.

Example:




What is Pair RDDs?

- Special form of RDD.
- Each element in an Pair RDD is a key-value pair.
- Key-value can be of any type.
- Common Pair RDDs are:
 - `groupByKey([numTasks])`
 - `reduceByKey(func, [numTasks]):`

Pair RDDs are useful:

- sorting, grouping, etc..



Creating Pair RDDs

- To create Pair RDDs:
 - get data in the form of key, value pair.
- Commonly used functions to create Pair RDDs:
 - groupByKey
 - reduceByKey

Pair RDDs – groupByKey([numTasks])

Objective: Display the list of IP addresses for each country .

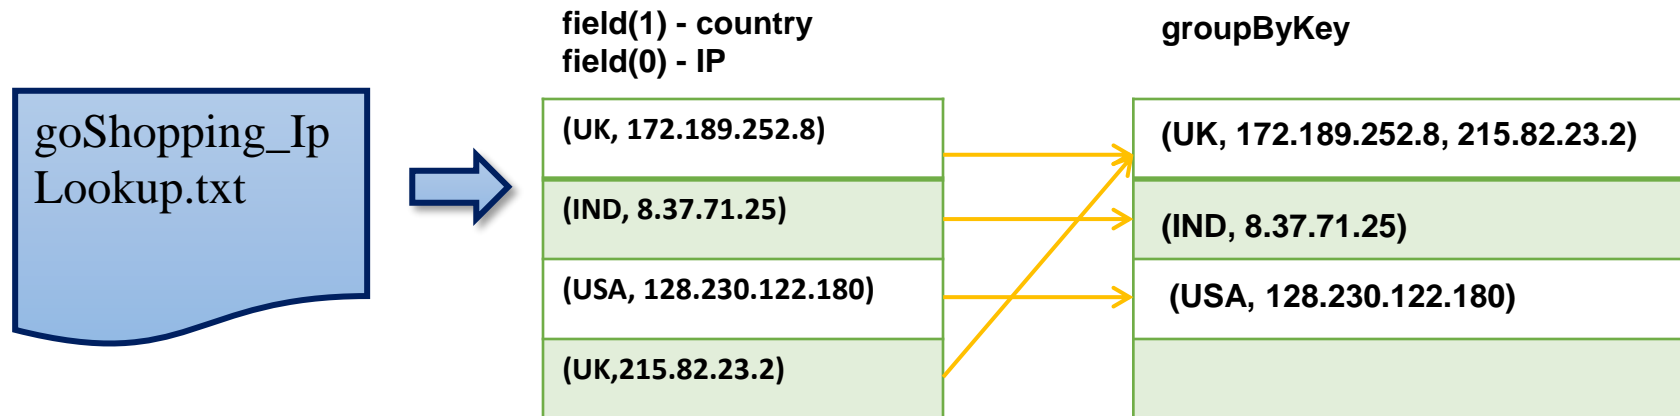
Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_IpLookup.txt")
val split = rdd.map(line => line.split(","))
val result = split.map(field => (field(1),field(0)))
result.groupByKey.take(10).foreach(print)
```

Output:

```
scala> result.groupByKey.take(10).foreach(print)
(UK,CompactBuffer(172.189.252.8, 215.82.23.2, 98.29.25.44, 68.199.40.156, 155.100.169.152, 38.68.15.223, 70.209.14.54, 74.111.6.173)) (IND,CompactBuffer(8.37.71.25, 8.37.71.69, 8.37.71.9, 8.37.71.57)) (USA,CompactBuffer(128.230.122.180, 128.122.140.238, 56.216.127.219, 54.114.107.209, 74.111.18.59, 8.37.70.170, 8.37.70.77, 8.37.70.112, 8.37.70.226, 8.37.70.99, 8.37.71.43))
```


Pair RDDs – groupByKey([numTasks])



Pair RDDs – reduceByKey(func, [numTasks])

Objective: Find total time spent on shopping site by each customer.

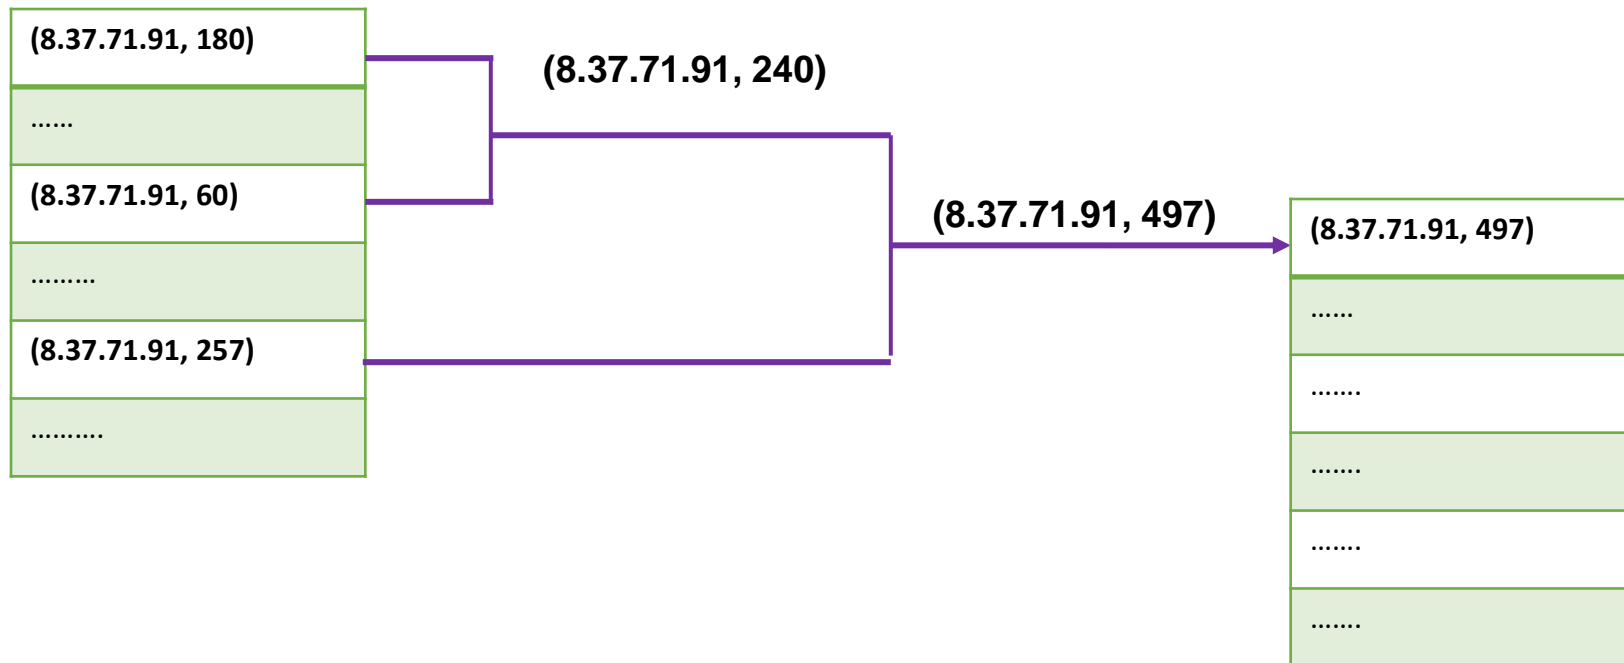
Action:

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")
val split = rdd.map(line => line.split("\t"))
val result = split.map(field => (field(4),field(6).toInt)). reduceByKey((v1,v2) =>
v1 + v2)
result.take(5).foreach(println)
```

Output:

```
(38.68.15.223,253943)
(56.216.127.219,203646)
(70.209.14.54,1406717)
(128.230.122.180,14112)
(54.114.107.209,559909)
```

Pair RDDs – reduceByKey(func, [numTasks])



Persisting RDDs

- One of the important features of Spark is is persisting (or caching) a dataset in memory across operations.
- Persisting an RDD:
 - each node stores any partitions of it that it computes in memory and reuses them in other actions on that dataset.
 - allows future actions to be much faster.
 - use `persist()` or `cache()` methods.
- The difference between `cache()` and `persist()`:
 - `cache()` == default storage level (`MEMORY_ONLY`).
 - `persist()` == specify various storage levels.

Storage Levels

Storage Level	Meaning
MEMORY_ONLY	<ul style="list-style-type: none">▪ Default level.▪ Store RDD as deserialized Java objects in the JVM
MEMORY_AND_DISK	<ul style="list-style-type: none">▪ Store RDD as deserialized Java objects in the JVM.▪ If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed.
DISK_ONLY	<ul style="list-style-type: none">▪ Store the RDD partitions only on disk.

Persisting RDDs - Example

```
val rdd = sc.textFile("/home/vagrant/dataset/goShopping_WebClicks.dat")  
import org.apache.spark.storage.StorageLevel  
rdd.persist(StorageLevel.MEMORY_ONLY)
```

Accumulators

Accumulators:

- Variables that are used for aggregating information across the executors.
- Simple syntax for aggregating values from worker nodes back to the driver program.

Use Case:

- Find out the number of blank logs (blank lines) .
- Number of times the network failed.
- Number of times zero sales were recorded.

Accumulators

Objective: Find the number of blank lines present in "goShopping_IpLookup.txt" .

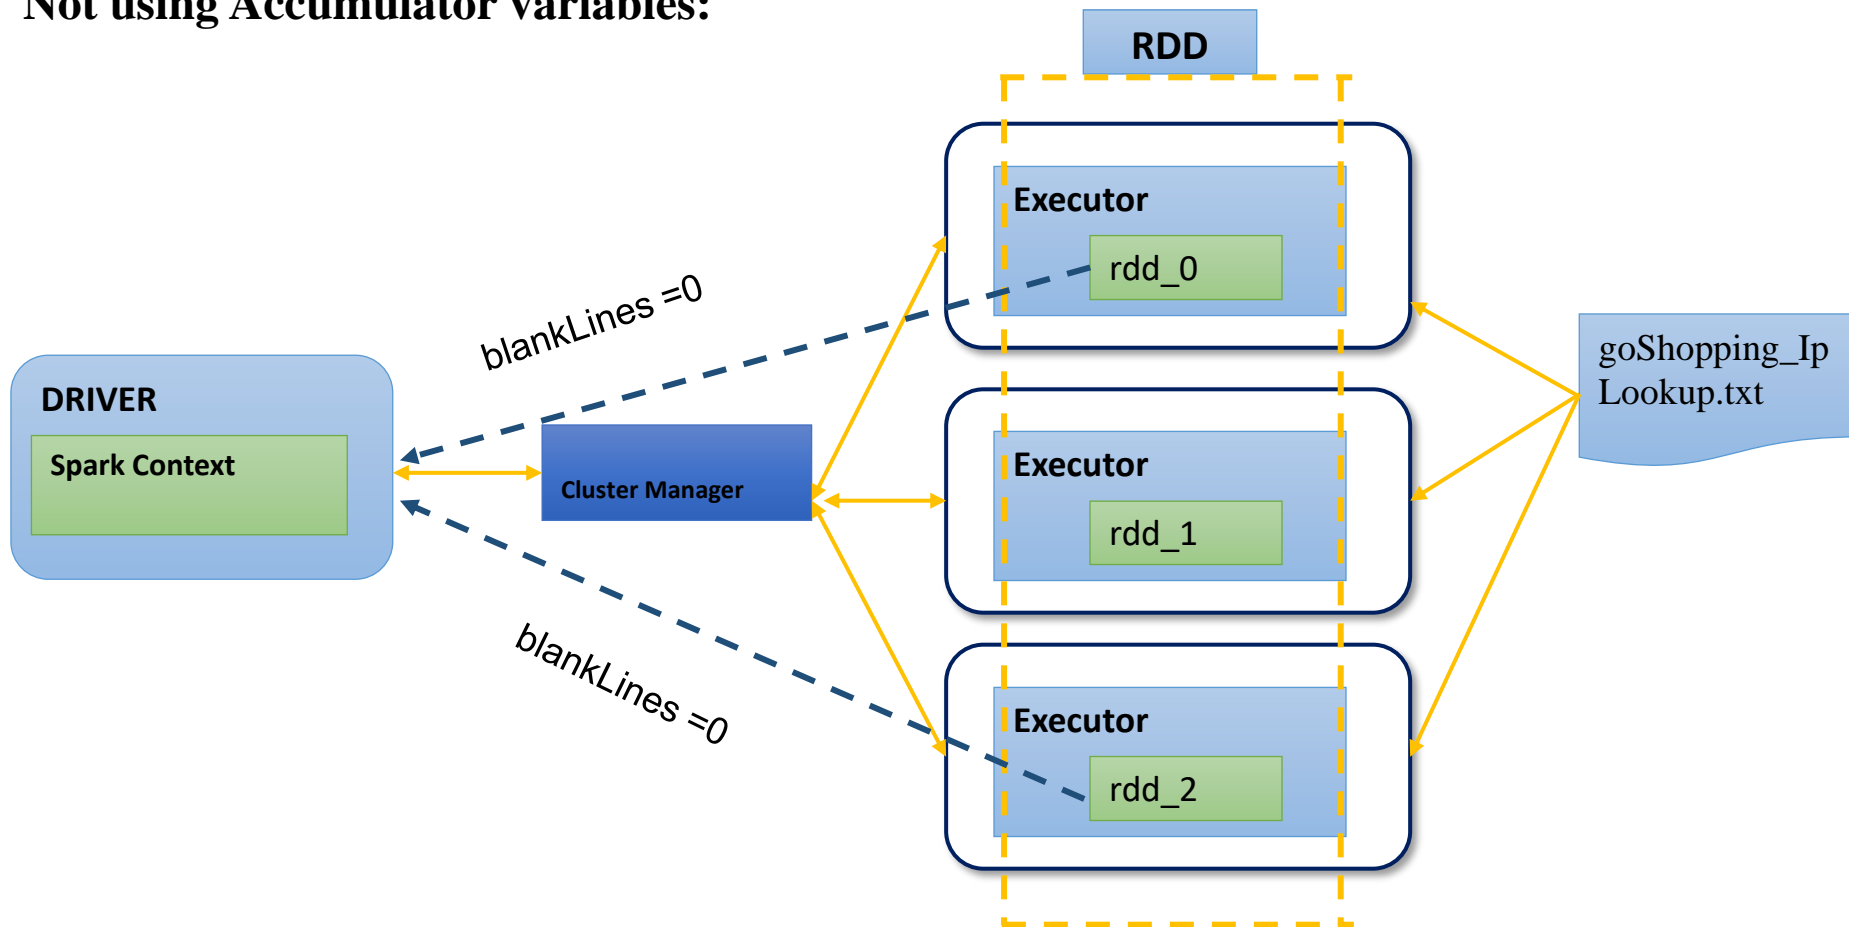
Input Data: "goShopping_IpLookup.txt"

```
172.189.252.8,UK,J5,Devon,38.955855,-77.447819  
215.82.23.2,UK,J6,Dorset,39.961176,-82.998794  
98.29.25.44,UK,J7,Down,41.49932,-81.694361
```

Number of blank lines: 2

Accumulators

Not using Accumulator variables:



Accumulators

Using Accumulator variables:

Action:

```
val blankLines = sc.accumulator(0)
sc.textFile("/home/vagrant/dataset/goShopping_IpLookup.txt").foreach{ line => if
(line.length() == 0) blankLines += 1 }
println("Blank lines: " + blankLines)
```

Output:

```
scala> val blankLines = sc.accumulator(0)
warning: there were two deprecation warnings; re-run with -deprecation for details
blankLines: org.apache.spark.Accumulator[Int] = 0

scala> sc.textFile("/home/vagrant/dataset/goShopping_IpLookup.txt").foreach{line => if (line.length() == 0)
  blankLines += 1}

scala> println("Blank lines: " + blankLines)
Blank lines: 2
```

Broadcast Variables

Broadcast Variables:

- To keep a read-only variable cached on each machine rather than shipping a copy of it with tasks.
- Use `SparkContext.broadcast(v)` to create broadcast variable.
- Variable `v` acts as a broadcast variable.
- `value()` can be used to access the broadcast variable.
- Useful when tasks across multiple stages need the same data.

Broadcast Variables

Objective:

Categorize the user IP as type "**AccidentVisit**" if time spent is less than 60 secs , else as type " **PurposeVisit** ".

" goShopping_WebClicks.dat "

Broadcast Variables

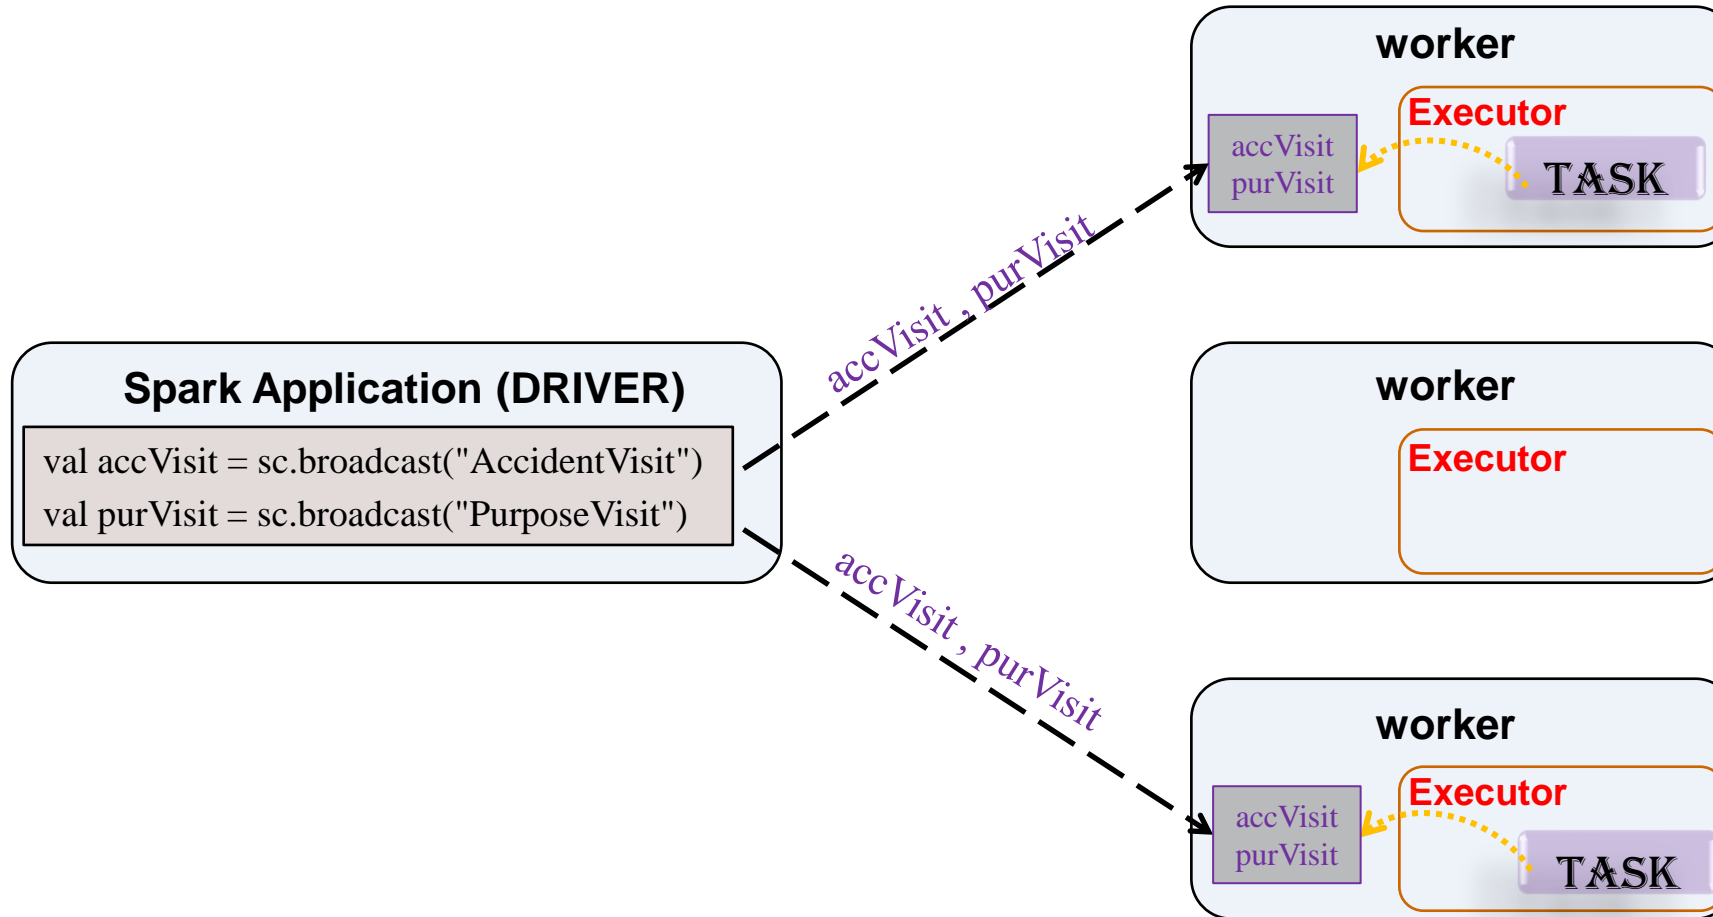
Action:

```
val accVisit = sc.broadcast("AccidentVisit")
val purVisit = sc.broadcast("PurposeVisit")
val rdd = sc.textFile("/home/vagrant/data/goShopping_WebClicks.dat",2)
val split = rdd.map(line => line.split("\t"))
val result = split.map(field =>
(field(4),if(field(6).toInt>=60){accVisit.value}else{purVisit.value}))
result.take(3).foreach(println)
```

Output:

```
scala> result.take(3).foreach(println)
(215.82.23.2,AccidentVisit)
(215.82.23.2,PurposeVisit)
(215.82.23.2,PurposeVisit)
```

Broadcast Variables



Questions



THANK YOU